

# Warhammer 40'000



In the grim darkness of the far future  
There is only WAR

# Why Machine Learning ?

A logical journey





# What is Warhammer 40'000 ?

## 1 – Turn-base Tabletop Game

- Played around a table
- Each player play alternatively

## 2 - Dark science-fantasy

- Set in the 41<sup>th</sup> millenium
- Every faction is evil...
- Or too stupid for that

## 3 – Next MCU

- The game / universe is more and more popular

# How is it played ?

## 1 – Army preparation

- Choose one of the 20+ factions
- Write down an army list within the agreed limit
- Collect and paint your models

## 2 – Meet your friend

- Reunite around a table
- Prepare the table setup
- Choose deployment and mission

## 3 – Play !

- Destroy your friend's army is not enough (*but still satisfying*)
- You must score more tactical objective points





# Game's Structure



# Movement Phase

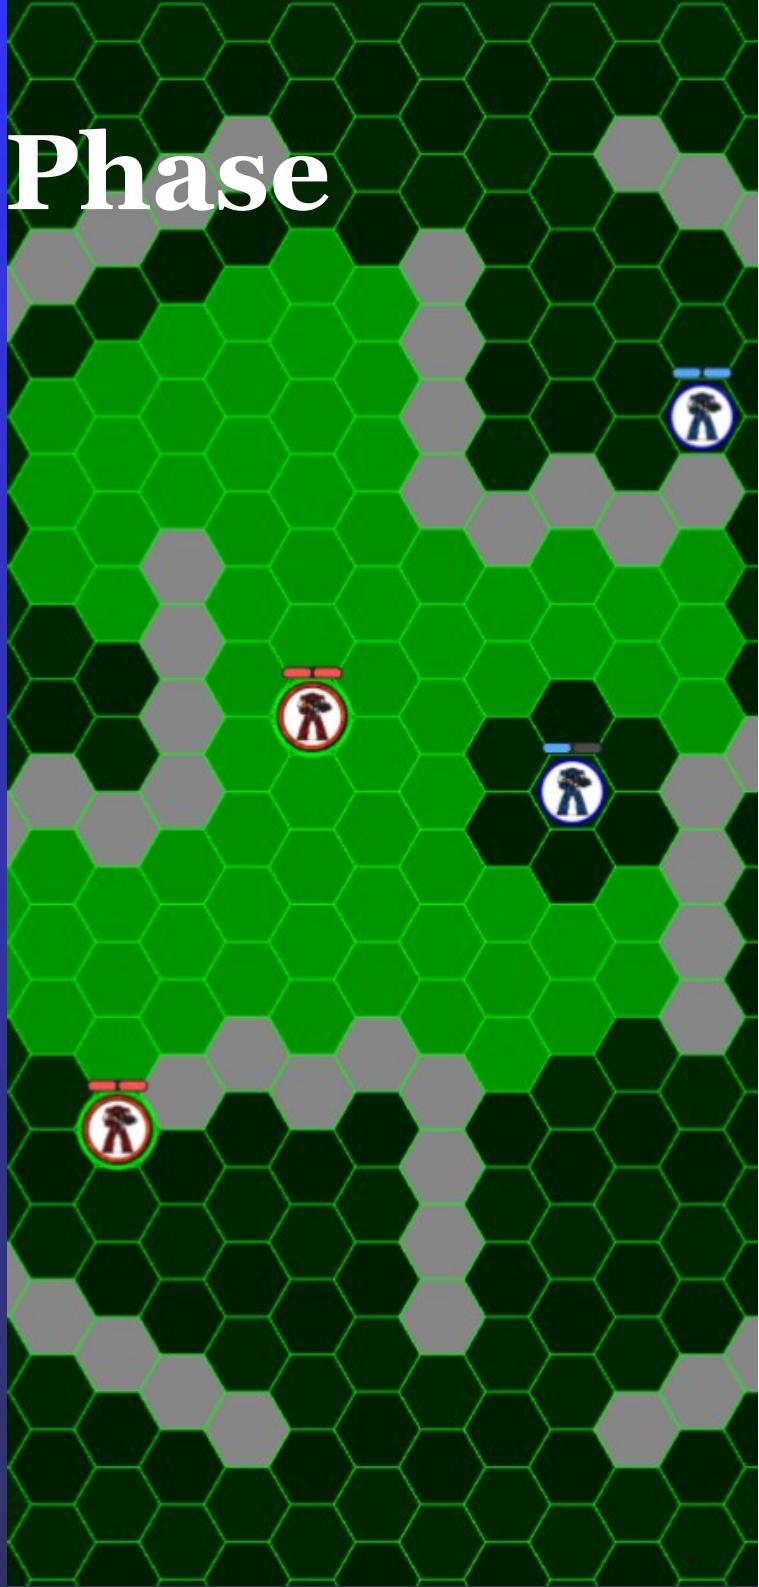
RANGE WEAPON									
M	T	SV	RNG	A	BS	S	AP	D	
6	4	3+	24"	2	3+	4	-1	1	

**Units can :**

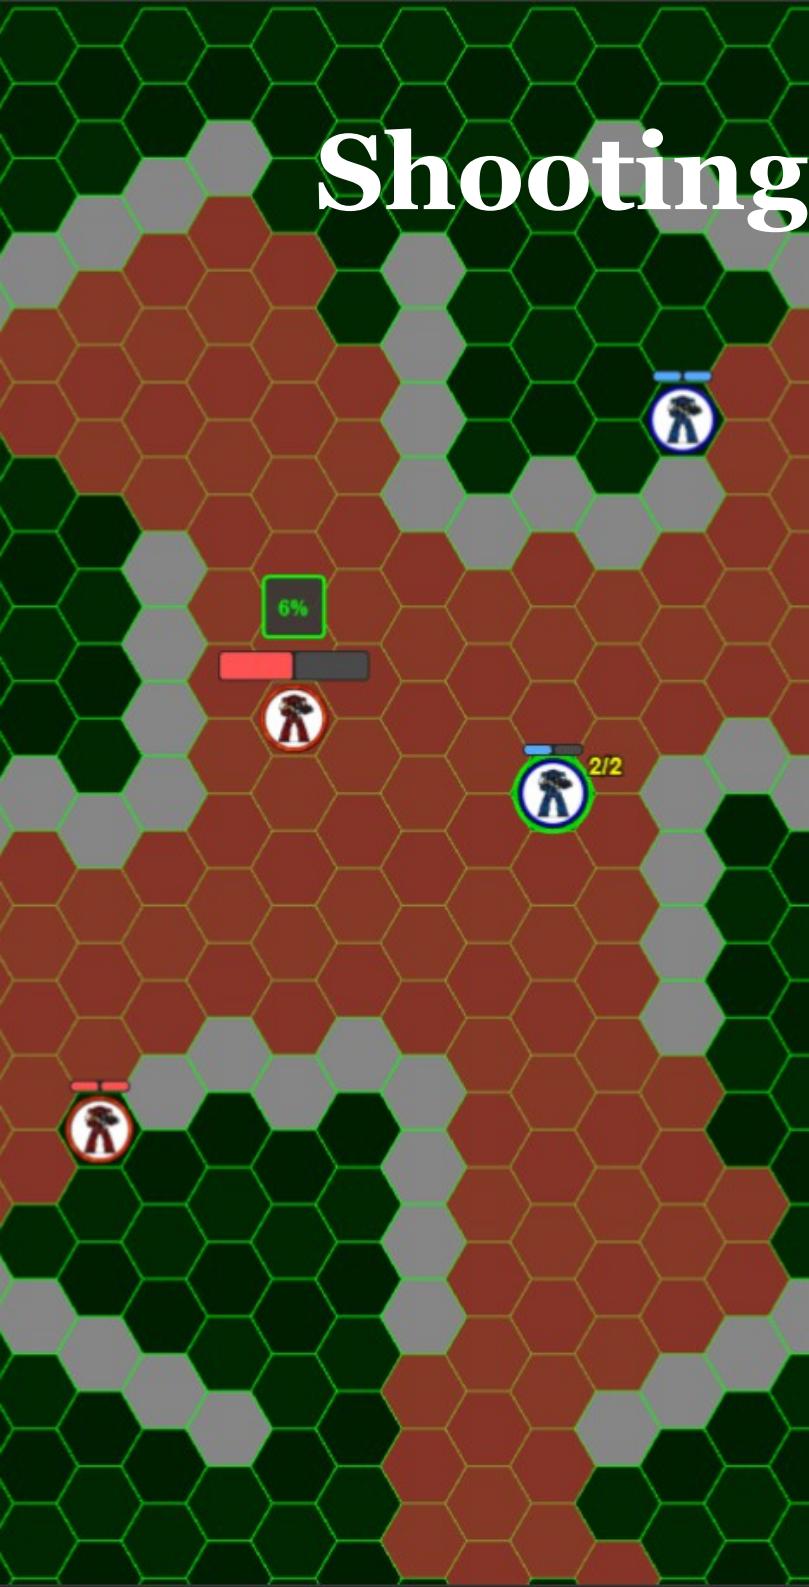
- Move up to **M** hexes

**Units cannot :**

- Move to or through a wall
- Move to another unit's hex
- Move to an hex adjacent to an enemy unit



# Shooting Phase



RANGE WEAPON								
M	T	SV	RNG	A	BS	S	AP	D
6	4	3+	24"	2	3+	4	-1	1

## Units can shoot :

- Up to their **A** times
- At units in Line of Sight
- At units within their **RNG**

## Shooting Sequence :

- Hit Roll : must roll  $\geq$  **BS**
  - ↳ Wound Roll : Weapon's **S** vs Target's **T**
    - ↳ Armor Save (**SV**) – **AP**
      - ↳ Inflict **D** damages

# The wargamer's doom

To play a 2 players game ...

**You need to be TWO !**





# You have a new friend !

No more need of an unfaithful human to play !

Welcome the Reinforcement Learning AI !!

A friend for life

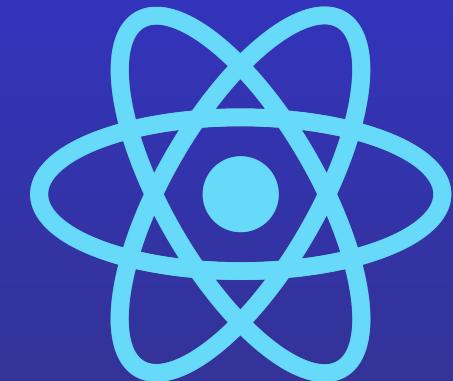
# Technical solution

*Make it real*

## Frontend : React

### **Game Requierments :**

- Board Management:
- Real-time UI Updates without full re-renders
- Event Handling: Click events



### **Performances :**

- PIXI.js Integration:
  - React manages PIXI.js canvas
  - TypeScript Compatibility

# Technical solution

*Make it real*

## Backend : Python



### **Obvious Choice :**

- Only viable option for serious deep reinforcement learning with DQN

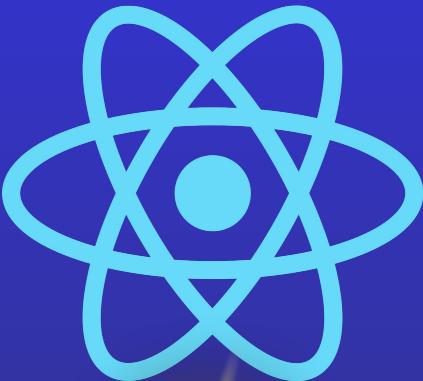
### **Familiarity :**

- I have made almost all the ML training on Python

# Technical solution

*The drama*

## Architecture conflict !

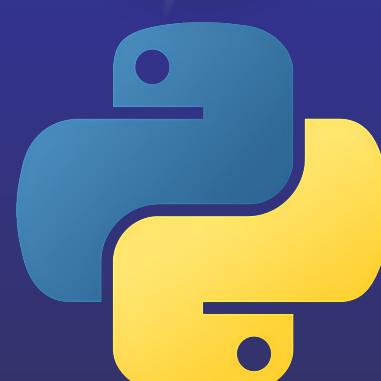


After 3 months of work, I had achieved :

- Complete PvP mode
- Fully functional agent training



**Time to play versus the AI !?**



I quickly realized a major flaw in my architecture :

- The PvP was led by React (click reactive)
- The AI was led by Python, obviously

**They could NOT work together !!!**

# Technical solution

*Mind your own Business*

## Architecture Alignment



### **Backend drives :**

- Sequential activation
- Phase transitions
- Game state authority
- AI decision making
- Action validation

### **Frontend displays :**

- Visual representation
- User input
- Animation/UX

### **Clear Boundary :**

**Backend drives ALL game logic → Frontend focuses on the display**

```

For each PLAYER unit
  ELIGIBILITY CHECK (Pool Building Phase)
    |- unit.HP_CUR > 0?
      |- NO → ✗ Dead unit (Skip, no log)
      |- unit.player === current_player?
        |- NO → ✗ Wrong player (Skip, no log)
    units_fled.includes(unit.id)?
      |- YES → ✗ Fled unit (Skip, no log)
    Adjacent to enemy unit within CC_RNG?
      |- YES → ✗ In fight (Skip, no log)
    unit.RNG_NB > 0?
      |- NO → ✗ No ranged weapon (Skip, no log)
    Has LOS to enemies within RNG_RNG?
      |- NO → ✗ No valid targets (Skip, no log)
    ALL conditions met → ✓ Add to shoot_activation_pool → Highlight the unit with a green circle around its icon
  STEP : UNIT_ACTIVABLE_CHECK → is shoot_activation_pool NOT empty ?
    YES → Current player is an AI player ?
      YES → pick one unit in shoot_activation_pool
        |- Clear any unit remaining in valid_target_pool
        |- Clear TOTAL_ATTACK log
        SHOOT_LEFT = RNG_NB
        While SHOOT_LEFT > 0
          |- Build valid_target_pool : All enemies within range AND in Line of Sight AND having HP_CUR > 0 → added to valid_target_pool
          valid_target_pool NOT empty ?
            YES → SHOOTING PHASE ACTIONS AVAILABLE
              |- Display the shooting preview (all the hexes with LoS and RNG_RNG are red)
              |- Display the HP bar blinking animation for every unit in valid_target_pool
              |- ✅ VALID ACTIONS: [shoot, wait]
              |- ✗ INVALID ACTIONS: [move, charge, attack] → end_activation (ERROR, 0, PASS, SHOOTING)
              AGENT ACTION SELECTION → Choose shoot?
                YES → ✅ VALID → Execute shoot
                  Agent choose a target in valid_target_pool
                    |- Execute attack_sequence(RNG)
                    SHOOT_LEFT -= 1
                    Concatenate Return to TOTAL_ACTION log
                    selected_target dies → Remove from valid_target_pool, continue
                    selected_target survives → Continue
                    GO TO STEP : PLAYER_ACTION_SELECTION
                    end_activation (ACTION, 1, SHOOTING, SHOOTING)
                NO → Agent chooses: wait?
                  YES → ✅ VALID → Execute wait action
                    |- end_activation (WAIT, 1, PASS, SHOOTING)
                  NO → Agent chooses invalid action (move/shoot/attack)?
                    |- ✗ INVALID ACTION ERROR → end_activation (ERROR, 0, PASS, SHOOTING)
                NO → end_activation (PASS, 0, PASS, SHOOTING)
            NO → Human player + STEP : UNIT_ACTIVATION → player activate one unit from shoot_activation_pool by left clicking on it
              |- Clear any unit remaining in valid_target_pool
              |- Clear TOTAL_ATTACK log
              SHOOT_LEFT = RNG_NB
              While SHOOT_LEFT > 0
                |- Build valid_target_pool : All enemies within range AND in Line of Sight AND having HP_CUR > 0 → added to valid_target_pool
                valid_target_pool NOT empty ?
                  YES → SHOOTING PHASE ACTIONS AVAILABLE
                    |- STEP : PLAYER_ACTION_SELECTION
                    |- Display the shooting preview (all the hexes with LoS and RNG_RNG are red)
                    |- Display the HP bar blinking animation for every unit in valid_target_pool
                      |- Left click on a target in valid_target_pool
                        |- Execute attack_sequence(RNG)
                        SHOOT_LEFT -= 1
                        Concatenate Return to TOTAL_ACTION log
                        selected_target dies → Remove from valid_target_pool, continue
                        selected_target survives → Continue
                        GO TO STEP : PLAYER_ACTION_SELECTION
                      |- Left click on another unit in shoot_activation_pool ?
                        SHOOT_LEFT = RNG_NB ?
                          YES → Postpone the shooting phase for this unit
                            |- GO TO STEP : UNIT_ACTIVABLE_CHECK
                          NO → The unit must end its activation when started
                            |- GO TO STEP : PLAYER_ACTION_SELECTION
                      |- Left click on the active_unit → No effect
                      Right click on the active_unit
                        |- SHOOT_LEFT = RNG_NB ?
                          NO → end_activation (ACTION, 1, SHOOTING, SHOOTING)
                            |- GO TO STEP : UNIT_ACTIVABLE_CHECK
                          YES → end_activation (WAIT, 1, PASS, SHOOTING)
                            |- GO TO STEP : UNIT_ACTIVABLE_CHECK
                        Left OR Right click anywhere else on the board
                          |- GO TO STEP : PLAYER_ACTION_SELECTION
                    NO → SHOOT_LEFT = RNG_NB ?
                      NO → shot the last target available in valid_target_pool → end_activation (ACTION, 1, SHOOTING, SHOOTING)
                      YES → no target available in valid_target_pool at activation → no shoot → end_activation (PASS, 1, PASS, SHOOTING)
                End of shooting → end_activation (ACTION, 1, SHOOTING, SHOOTING)
              End of shooting phase → Advance to charge phase

```

# Preparation Work

*Make the game follow ONE ruler*

This « algorythm » describes the way the shooting phase should be played

# Back to the AI





# What is Q learning ?

*Overview*

You are playing a new video game... like a kid!



Over time, you learn what works and what doesn't

**This is the way Q-Learning works**

# How to help Bob to survive ?

*Q learning to the rescue !*

## Bob's plan to survive :

### 1 - Bob explores his environment

→ "What if I charge a Bloodthirster or rush an objective ?"

### 2 – Bob builds a strategy

→ "A Bloodthirster is too dangerous, go for the objective"

Over time, Bob remembers what works best and makes better choices !

(think Edge of tomorrow)





# The Formula

$$Q(s,a) = Q(s,a) + \alpha [r + \gamma \max Q(s',a') - Q(s,a)]$$

$Q(s, a)$  → The current guess for how good the action ( $a$ ) in this state ( $s$ ) is

$\alpha$  → How fast he updates what he knows

$r$  → Reward

$\gamma$  → How much we value current over future rewards

$\max Q(s', a')$  → The best future reward he can expect

This function will return how good is it to be in the state **s** and take an action **a** in this state

# Learning parameters

## Learning Rate ( $\alpha$ ) : How fast Bob adapts?

- Learns too fast → Forms an opinion too quickly
- Learns too slow → Keeps repeating the same mistakes

✓ Smart Bob :  
Remembers key lessons but still adapts

## Future Discount ( $\gamma$ ) : Thinking Long or Short-Term?

- Short-term → Misses big rewards
- Long-term → Waits too long & loses

✓ Smart Bob :  
Balances small wins & big rewards



This is a Bloothirster

# The Q table

Think of the Q-table as Bob's personal cheat sheet



<b>Q</b>	<b>a1</b>	<b>a2</b>	<b>a3</b>	<b>a4</b>
<b>S1</b>	$Q(S1, a1)$	$Q(S1, a2)$	$Q(S1, a3)$	$Q(S1, a4)$
<b>S2</b>	$Q(S2, a1)$	$Q(S2, a2)$	$Q(S2, a3)$	$Q(S2, a4)$
<b>S3</b>	$Q(S3, a1)$	$Q(S3, a2)$	$Q(S3, a3)$	$Q(S3, a4)$
.	.	.	.	.

**Goal :**  
Update and learn these  
Q-values such as the  
total reward is  
maximized

A detailed Warhammer 40k illustration depicting a confrontation between two Space Marine chapters in a vast, ornate hall. On the left, a chapter of Space Marines in dark red and black armor, featuring a prominent eye symbol on their chest plates, stands in formation. On the right, another chapter in gold and red armor, with a large golden eagle emblem on their shoulder pads, also stands in formation. Between them is a large, glowing orange pyramid-shaped object on a platform, with several small figures standing around it. The background shows the grand architecture of the hall, with high ceilings, decorative columns, and warm lighting from hanging lanterns.

Apply it to Warhammer 40k



# Observation

How Bob knows his environment ?

## DQN Input

### 142 positions array :

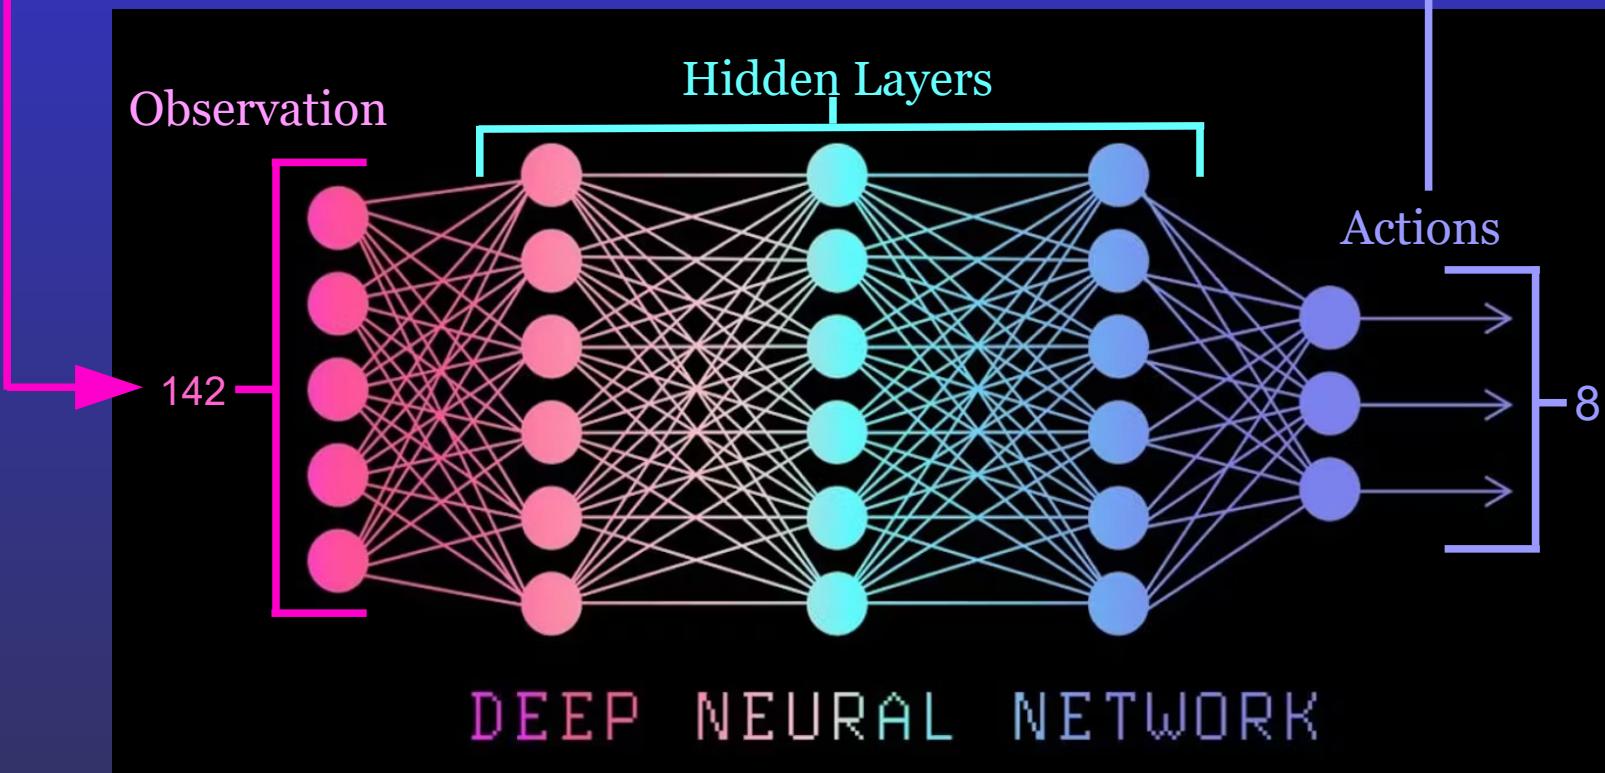
- General game status (Active player, turn, etc...)
- Informations about all the units on the board
- Situation on the board
- Board structure

# From Q-table to DQN

How does Bob handles his environment ?

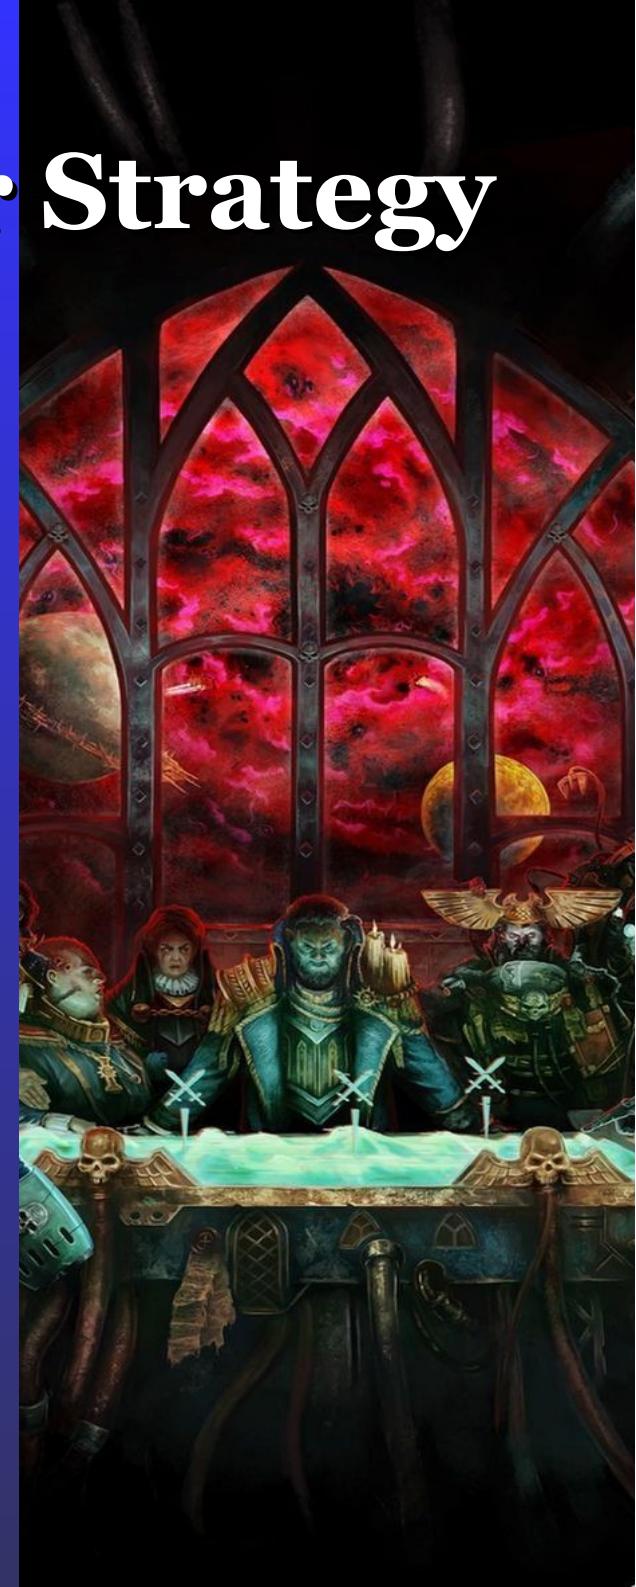
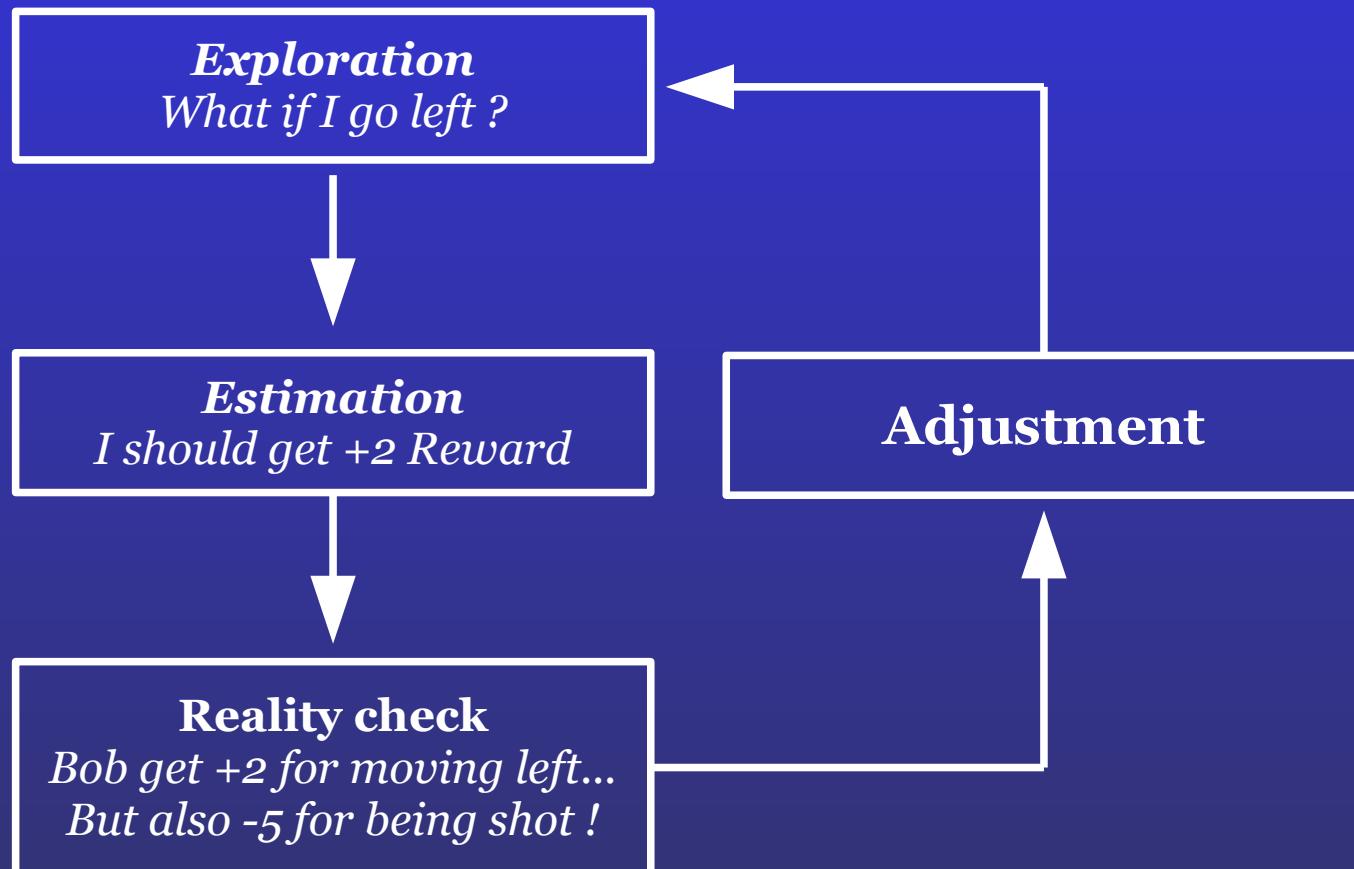
(Enter Bob's brain)

Q	a1	a2	a3	a4
S1	$Q(S1, a1)$	$Q(S1, a2)$	$Q(S1, a3)$	$Q(S1, a4)$
S2	$Q(S2, a1)$	$Q(S2, a2)$	$Q(S2, a3)$	$Q(S2, a4)$
S3	$Q(S3, a1)$	$Q(S3, a2)$	$Q(S3, a3)$	$Q(S3, a4)$
.	.	.	.	.



# Tailor Strategy

How Bob learns to play ?





# Rewards

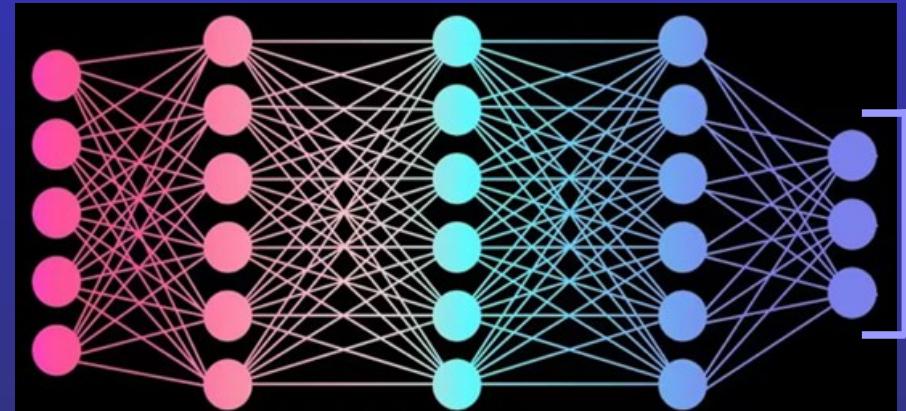
How Bob knows if he is doing good ?

```
"SpaceMarine_Infantry_Troop_RangedSwarm": {  
    "base_actions": {  
        "move_close": 0.4,  
        "move_away": 0.5,  
        "move_to_charge": 0.2,  
        "move_to_los": 0.6,  
        "ranged_attack": 0.5,  
        "melee_attack": 0.1,  
        "charge": 0.1,  
        "charge_success": 0.1,  
        "being_charged": -0.3,  
        "wait": -0.5  
    },  
    "target_type_bonuses": {  
        "vs_melee": 0.1,  
        "vs_ranged": 0.0,  
        "vs_swarm": 0.2,  
        "vs_troop": 0.1,  
        "vs_elite": 0.0,  
        "vs_vehicle": -0.1,  
        "vs_heavy": -0.2  
    },  
    "result_bonuses": {  
        "kill_target": 0.2,  
        "wound_target": 0.1,  
        "no_overkill": 0.1,  
        "target_lowest_hp": 0.1  
    }  
}
```

# Make a Choice

What will Bob decide to do ?

Action taken : Action with the highest Q-value



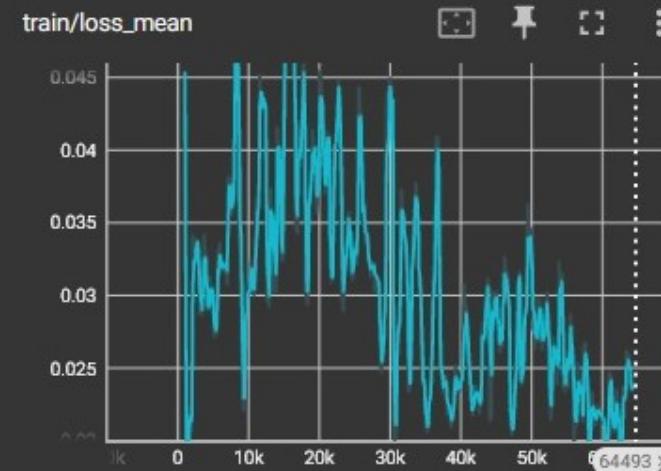
```
action_map = {  
    0: "move_north",  
    1: "move_south",  
    2: "move_east",  
    3: "move_west",  
    4: "shoot",  
    5: "charge",  
    6: "attack",  
    7: "wait"  
}
```



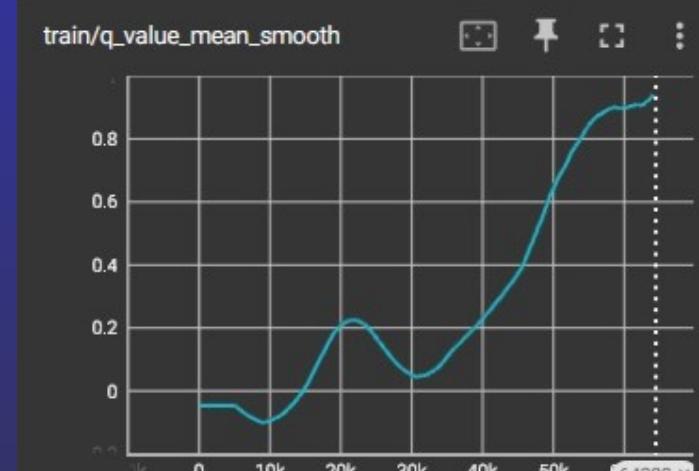
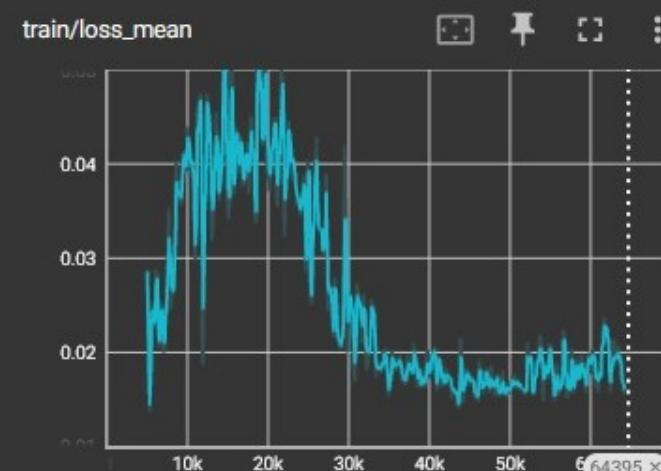
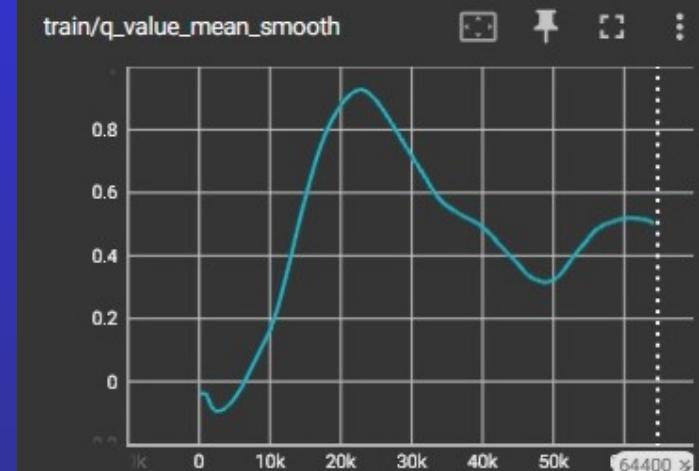
# Supervision

How to ensure Bob is well trained ?

**Loss**



**Q-Value**



# Let's Playyyy !!!



# Future improvements

## 1 - Finish the game's core

- Add the charge and combat phases

## 2 – Upgrade Observation

- Add tactical datas

## 3 – Optimize the agent training

- Make it worth playing against

## 4 - Add new units

- Assault Intercessors, Captain Gravis....

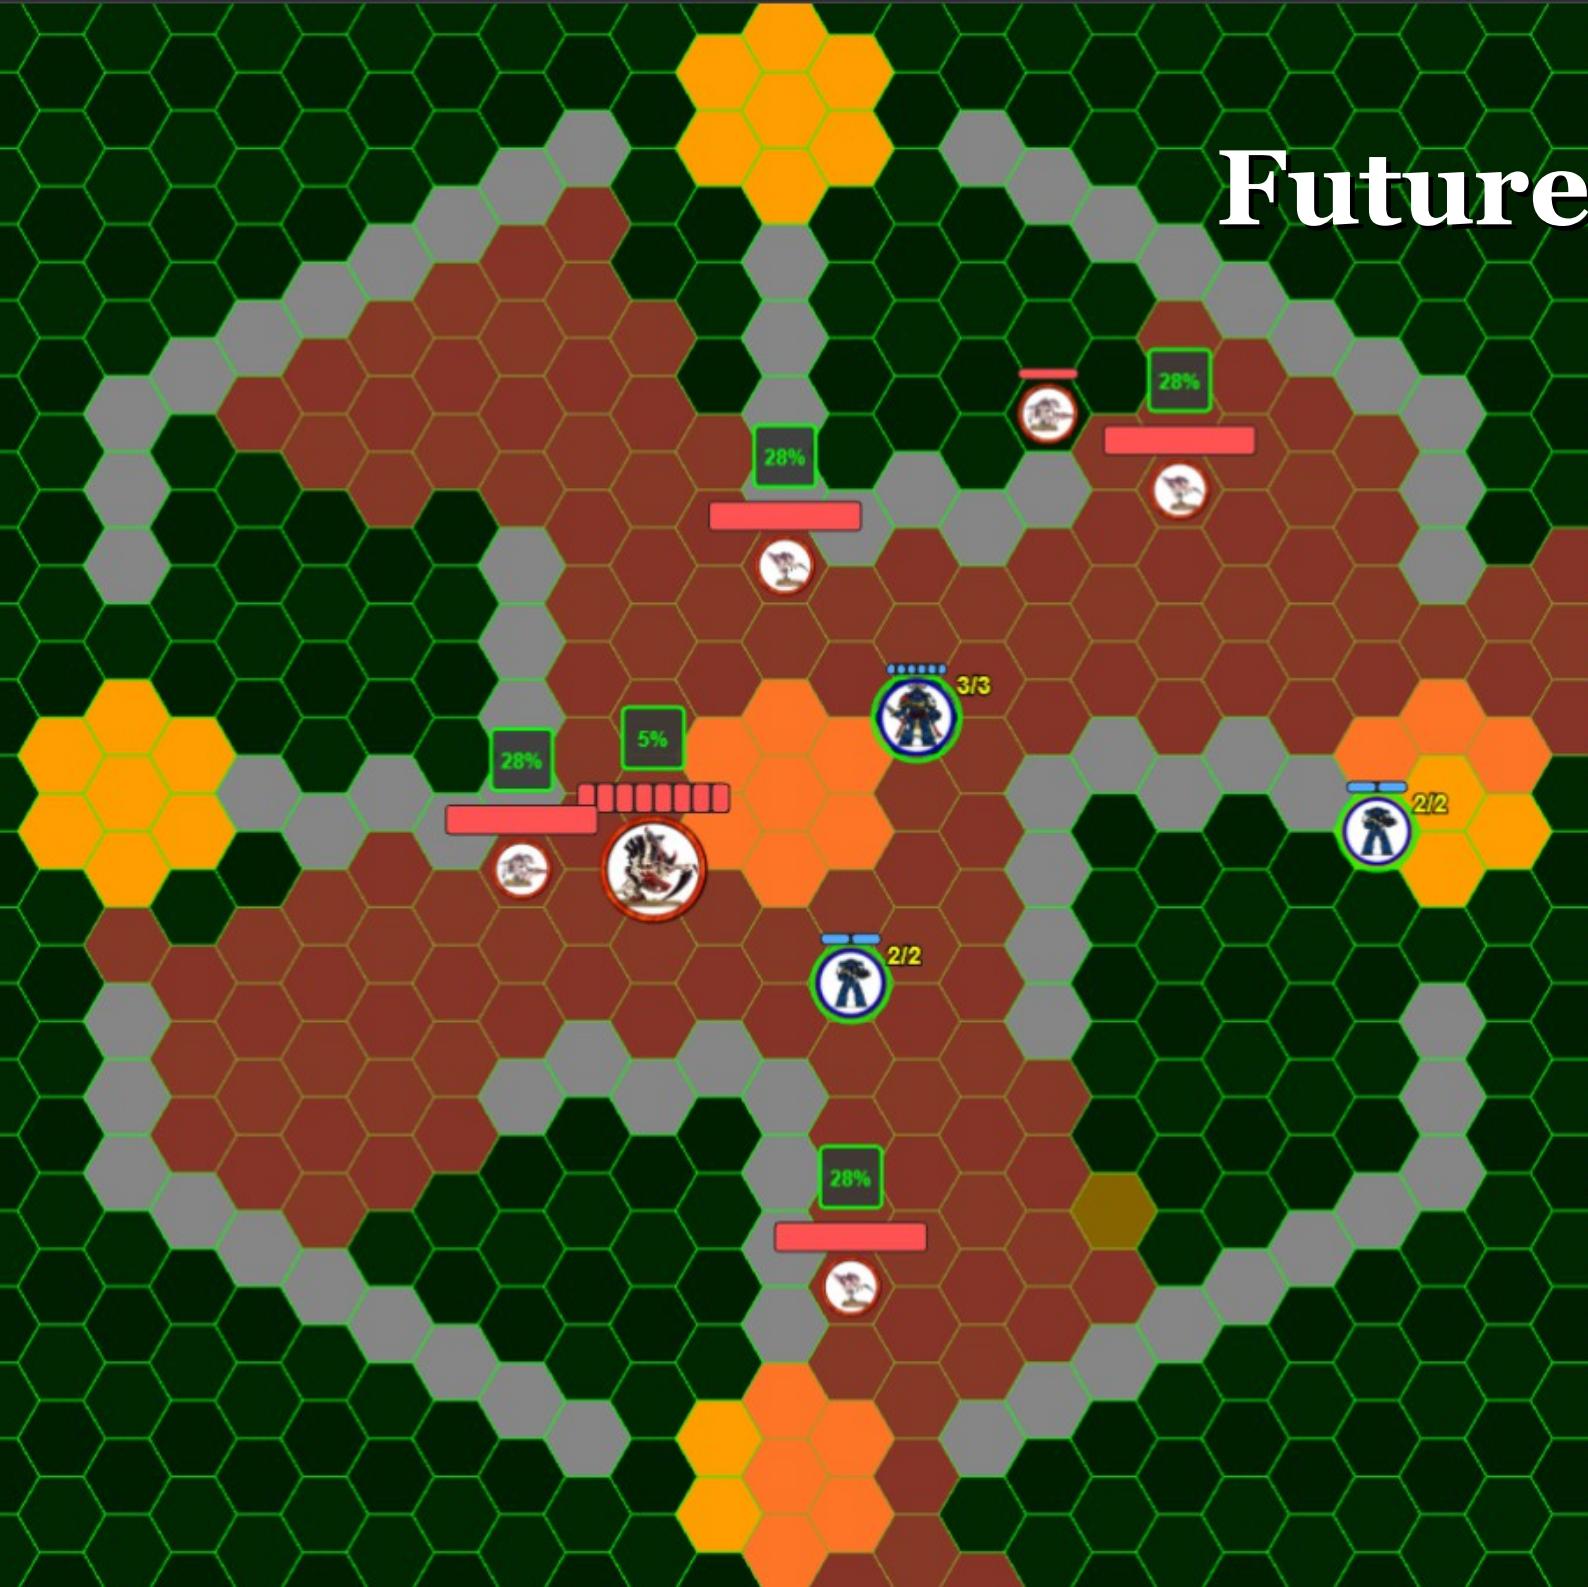
## 5 - Add new factions

- Tyranids, Orks, etc...

## 6 - Share the game !



# Future Features



## Complete the turn

- Add the last phases

## Tactical Objectives

## New Marins Unit

Captain Gravis

## New Faction

Tyranids

## Tyranids Unit

Carnifex  
Hormagaunt  
Termagant

# My Future

Find a first job in AI

...

Anywhere in the galaxy !

(and keep learning)

